

## CA169: Week 2

### Encoding, Transmission Errors and Protocols

#### Recap from week 1


- Last week we looked at the physical ways we can transmit data
  - Unshielded Twisted pair (Phone & Ethernet)
  - Coaxial (Aerial TV, fibre to home)
  - Fibre optics (Fibre internet)
  - Radio
    - Bluetooth
    - Wifi
    - NFC

#### Transmission & Encoding


- How do we get 1's and 0's from electrical signals?
  - Encoding
- How do we deal with errors from these signals ?
  - Checksums

#### Analog v Digital

- Physical cables deal with voltage
- Digital systems prefer binary (1's and 0's)
- Our goal is to turn a signal from a line (analog) into a signal a computer can understand (digital)
- Essentially, turn this:



- Into this:



- This problem is harder than it seems
- What about a weak signal




- Or a signal with a lot of interference?




- We use an encoding algorithm
- There is no one size fits all solution
  - But some are better than others
- There are a number of factors to consider when choosing one
  - Physical medium (ethernet, wifi etc...)
  - Distance
  - voltage

### Algorithm 1: NRZ-L & NRZ - I

- Non Return Zero Level (NRZ-L) is an encoding scheme where high voltage becomes a 1 and low voltage becomes 0
- This is used for short connections (e.g. to a printer)




- Non Return Zero Inverted (NRZ-I) doesn't work off high-low voltage
- It looks at the transitions between positive and negative voltage
- If there is a transition, we use 1
- If there is no transition, we use 0



- Which to choose?
  - NRZ-I can deal with noise better
  - It is easier to check if voltage went from - to +
  - Than to say "if voltage >5"
  - Problem
  - What if I wanted to send the data 11111111

### Clock Information

- So we clearly have to introduce a time component (clock)
- Manchester encoding does this
- We cut our signal into segments (based on time)
- We call this Biphase
- If we go from + to - we use a 0
- If we go from - to + we use a 1



### Why Manchester Encoding

- More difficult to operate
- More resistant to noise
- Requires twice as much bandwidth to send the same data as other methods
- But much more reliable

## 4B/5B Coding

- 50% more efficient than Manchester (we use it in 100Mb internet)
- Take four bits from signal
- Encode them as 5 bits (we call this a cell)

Data		4B5B code
(Hex)	(Binary)	
0	0000	11110
1	0001	01001
2	0010	10100
3	0011	10101
4	0100	01010
5	0101	01011
6	0110	01110
7	0111	01111

- By mapping 4 bits onto 5 bits we have better control over the data
- This makes it more resistant to noise
- If there are only 4 bits, there are 16 possible frames ( $2^4$ )
- In short, we have a much more reliable signal

## Modems

- This process of converting from analog to digital (and back again) is called modulation
- In the real world we use MODEMs to do this
- MODEM is short for:
  - MODulator
  - DEModulator

## Transmission Errors

- We know how to convert our analog signal into a digital one
- But how do we know we converted it right?
- A number of factors can influence our transmission
- We call these noise
- Noise can come in many forms
  - Electrical interference
  - Walls
  - Long distance

## How to Tell if Data is Incorrect

- How can we tell if our data is incorrect?
- By adding extra check bits onto the end of our transmission
- And doing a bit of maths

## Cyclic Redundancy Code (CRC)

- Error detecting codes
- Not error correcting codes
- Idea: represent binary as a polynomial

$$F = 110001$$

$$F(x) = (x^5 * 1) + (x^4 * 1) + (x^3 * 0) + (x^2 * 0) + (x^1 * 0) + (x^0 * 1)$$

$$F(x) = x^5 + x^4 + 0 + 0 + 0 + x^0$$

$$F(x) = x^5 + x^4 + x^0$$

$$F(x) = x^5 + x^4 + x^0$$

## CRC Algorithm

- Compare using Modulo 2 arithmetic ( $x\%2$ )
  - This is the same as the XOR operation  $\oplus$
- Sender & receiver agree on  $G(x)$  generator polynomial.
- Append  $R$  0 bits to  $M(x)$ , the message, where  $R$  is the degree of  $G(x)$ , this yields
- $x^R * M(x)$
- Divide  $G(x)$  into  $x^R * M(x)$ .
- Add remainder to  $x^R * M(x)$ , result  $T(x)$
- Example

$M(x) = 1101\ 0110\ 11$

$x^R * M(x) = 1101\ 0110\ 1100\ 00$

$G(x) = 10011$  or  $x^4 + x^1 + x^0$


Remainder should be 1110

$T(x)$  transmitted message should be

1101 0110 11 1110

This will be evenly divisible by  $G(x)$  !

Frame : 1 1 0 1 0 1 1 0 1 1  
 Generator: 1 0 0 1 1  
 Message after appending 4 zero bits: 1 1 0 1 0 1 1 0 0 0 0



Transmitted frame: 1 1 0 1 0 1 1 0 1 1 1 1 0

- Standard CRCs
  - CRC-12
    - $x^{12} + x^{11} + x^3 + x^2 + x + 1$
  - CRC-16
    - $x^{16} + x^{15} + x^2 + 1$
  - CRC-CCITT
    - $x^{16} + x^{12} + x^5 + 1$
  - CRC-32
    - $x^{32} + x^{26} + x^{23} + x^{22} + x^{16} + x^{12} + x^{11} + x^{10} + x^8 + x^7 + x^5 + x^4 + x^2 + x + 1$

## How Data is Transmitted

- I want to send a document to another device
- This document is converted into 1's and 0's
- It is sent to the other device
- The other device then converts the 1's and 0's back into the document

## IRL

- In reality, this is much harder
- The document isn't converted into one long stream of 1's and 0's
- It is broken down into packets/frames of information
- These packets are then sent to the other device
- Each packet contains only one part of the document
- Why do we do this?
  - If we sent the entire document in one go
  - We can't handle sending multiple things at once (multiplexing)
  - Your computer sends a lot of information at once all the time
  - See netstat
- What's the problem?
  - How do we tell when one packet/frame ends and another begins?
  - We call this the framing problem

## Possible Solutions to the Framing Problem

- Time delay:
  - Every 1 second send a new packet
  - Data may be lost if signal is interrupted, or a delay in transmission
- What if we include the length of the packet at the end
  - Add on another bit to the CRC
  - If the count is corrupted, we would never know how long our data is

## Actual Solution

- At the start and end of our packet we add in special markers
  - We call this character stuffing
- DLE STX (meaning start of packet)
  - Lets pretend this is: 1001 in binary
- DLE ETX (end of packet)
  - Lets pretend this is 0110 in binary

## Examples

- If I want to send 11010101 to another computer
  - We add on our DLE STX to the start
    - 100111010101
  - Then we add on our DLE ETX to the end
    - 1001110101010110
- The computer we are sending the data to knows that a start packet is 1001 and an end packet is 0110
- When it gets this data it removes them
- This leaves us with our original data

### Double Stuffing Example

- If I want to send the data 110101011001010
- Find bits that match STX or ETX 110101011001010
- Duplicate them 1101010110011001010
- Now add on the STX and ETX to the packet
- 100111010101100110010100110
- If we see a duplicated pattern, only remove one of them!

### What Does Protocol Mean

- We now know how to change data from signals into nice 1's and 0's
- We can also be sure that the data we sent is correct
- But there are a number of questions we still have to answer
- What do we do if we do get an error?
- How do we signal that we have finished sending data?

### Need for Protocols

- Protocols dictate how we communicate
- They control
  - Data format
  - Data content Timing/synchronisation of communication
  - How the communication is managed

### Why Use Protocols

- In networking a protocol is like a language
- It provides a common way of devices to understand each other
- There is no magic
- They are designed by programmers
- They are only useful if they are widely adopted
- They are refined through trial and error

### Making a Protocol

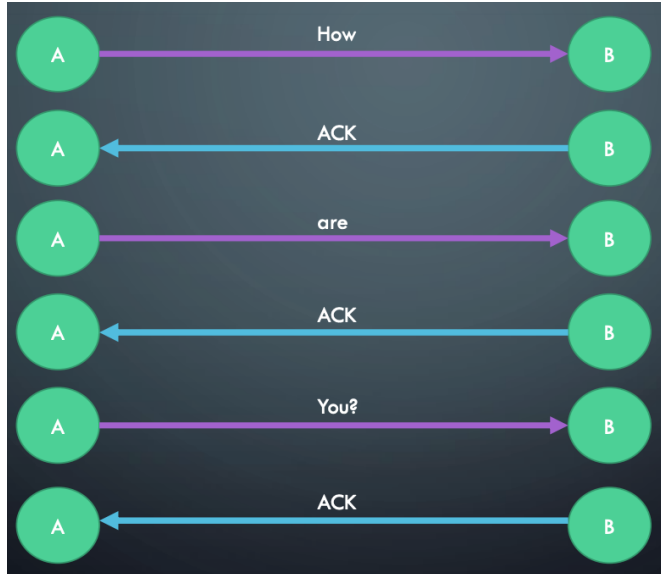
- Lets first make some assumptions:
  - Transmission will be simplex (only one computer can speak at a time)
  - Computers are always ready to receive data
  - Processing time is negligible and we have unlimited space
  - Our physical medium is error free!

### Utopia Protocol

- Make all assumptions, a, b, c, d.
- For this protocol the control header and checksum are unnecessary.
- The transmitting host simply takes packets from host A (which always has one ready) and pumps them as fast as it can onto the physical link.
- The receiver accepts the frames and passes them straight to host B

## Stop & Wait

- The UTOPIA protocol is nice but...
- Processing time is not negligible, and we do not have infinite space
- So we drop assumption c
- In practice the receiving computer needs time to process the packets
- Only has a fixed amount of space to store incoming packets for processing
- The receiver needs a way to prevent it from getting more data than it can handle
- When the receiver processes a packet, it sends one back requesting more data (we'll call this a control frame)
- Only after receiving this control frame will the sender fetch and transmit the next packet
- E.g Let's send the data "How are you"?



- Issues?
  - Can anyone see any problem with this method?
  - What about a failed transmission?

## Positive Acknowledgement With Retransmission (PAR)

- Noisy channel simplex protocol: drop some assumptions
  - c) [processing time, buffers pace] and
  - d) [error free link]
- With error prone physical link, frames may be either damaged or lost completely.
- However,
  - Damaged frames can be detected by the checksum.
  - Lost frames will not be acknowledged.
  - Eventually the sender will tire of waiting for an acknowledgement, timeout, and retransmit the frame.
- Same as Stop-and-Wait, except that
  - damaged frames are not acknowledged,
  - causing a timeout and subsequent retransmission

## Examining Another Protocol

- Suppose the message "Get Lost Fred" is being sent from A to B, one word per packet



- What happened?
- B receives "Get lost Fred Fred" and the protocol has failed !
- Basically what has happened is that the receiver has accepted a duplicate frame.
- The solution is to use SEQ, a sequence number in the control header to differentiate between frames and allow duplicates to be discarded.
- As a frame must be acknowledged before the next one is sent, a one-bit sequence number [0,1] is sufficient.
- The receiver will expect alternatively numbered frames (0 1 0 1 0 1 ... etc.).
- Any frame with the wrong sequence number is rejected as a duplicate
  - (but still acknowledged).



- What about time outs





## Timeouts

- The message "Get Lost" is received and the protocol has failed

## Solution

- Include in the ACK field of the acknowledgement control frame the SEQ number of the last frame received without error.
- Then if this number (0 or 1) differs from the transmitted frame, the sender transmits it again.
- The resulting PAR simplex protocol will now work in the face of any combination of
  - garbled frames,
  - lost frames and
  - premature timeouts.



## Summary

- After transmitting a frame and starting the timer, Host A waits for a response.
- There are three possibilities:
  - 1. an acknowledgement frame arrives undamaged
  - 2. a damaged acknowledgement arrives, or
  - 3. the timer goes off.
- If a valid ACK comes in, A fetches the next packet and puts it in the buffer
  - overwriting the previous packet, and advancing the sequence number.
- If a damaged frame arrives or no frame at all arrives,
  - neither a buffer nor the sequence number are changed, so that a duplicate can be sent

## Bi-Directional PAR

- The control header of all frames contains the three fields KIND, SEQ and ACK.
- Sending data packets/acknowledgements in both directions is no problem –
  - by looking at the KIND bit in the header, the receiver knows which it is dealing with.
- However, this would be inefficient.
- Consider an Host B which is about to acknowledge a data frame received from
- Host A, and also about to send off a data frame to A.
- Instead of sending two frames, the acknowledgement can hitch a lift on the data frame, using the ACK field in the header.
- This is called piggybacking.
- As we are still making assumption b)
- [Transmitting and receiving hosts always ready to transmit & receive data]
- All acknowledgements can be piggybacked. Thus, data packets are bounced back and forth between A and B.
- Notes:
  - For the protocols considered so far, only one frame is in the wire at any one time.
  - The sender needs to keep a copy of each frame in a buffer for possible retransmission until the frame has been successfully acknowledged.
- Assumption b)
- [Transmitting and receiving hosts always ready to transmit & receive data]
- Easily dealt with.
  - If there is no outgoing data frame, the host will wait a short while to see if one comes along to provide a piggyback.
  - If not, a separate acknowledgement frame will be sent.
  - It must not wait too long to avoid unnecessary duplicates being sent due to the sender timing-out.
- Up until now, lost and damaged frames have been dealt with in the same way.
- No ACK is sent, leading to timeout and retransmission.
- The timeout period is usually set quite long in order to avoid complications caused by premature timeout.
- This is inefficient, as while the timeout is expiring, the link is not being used.
- A partial solution is NAK, a negative acknowledgement.
- This is sent immediately a damaged frame is received and elicits immediate retransmission.
- The NAK may also be piggybacked.
- If the NAK is damaged, no harm is done as the sender will eventually timeout and retransmit as before.
- A damaged frame is Nak'd only once.

## Pipelining

- When propagation delay is not negligible, these previous methods are wasteful of bandwidth.
- The solution is to 'fill up the pipe'.
- However, doing this entails sending off frames before ACKs for previous frames have arrived.

## Sliding Window Protocol

- Each outbound frame is given a sequence number in the range of  $2^n - 1$  using an n-bit field, e.g. if  $n=1$ , then range is 0.....1 as in ABP or PAR protocols.
- Both sender and receiver keep windows informing them of which frames can be validly sent and which validly received.

## Rules

- Sender :- The upper edge of sender is advanced when a frame is sent (up to max. window size).
- The lower edge advanced when ACK received for lowest numbered frame in the window.
- Receiver :- Both edges are advanced when the lowest numbered frame in window is correctly received and ACK sent.

## Notes

- Buffering requirements at both sender and receiver depend on the size of the sending and receiving windows respectively.
- Each transmitted frame has its own separate timeout clock.
- In these protocols, an acknowledgement for frame N is accepted as acknowledging all transmitted frames numbered up to N (counting circularly).
- Thus, if ACK(0) and ACK(1) were both destroyed, but ACK(2) now arrives, it implicitly acknowledges 0 and 1 also.


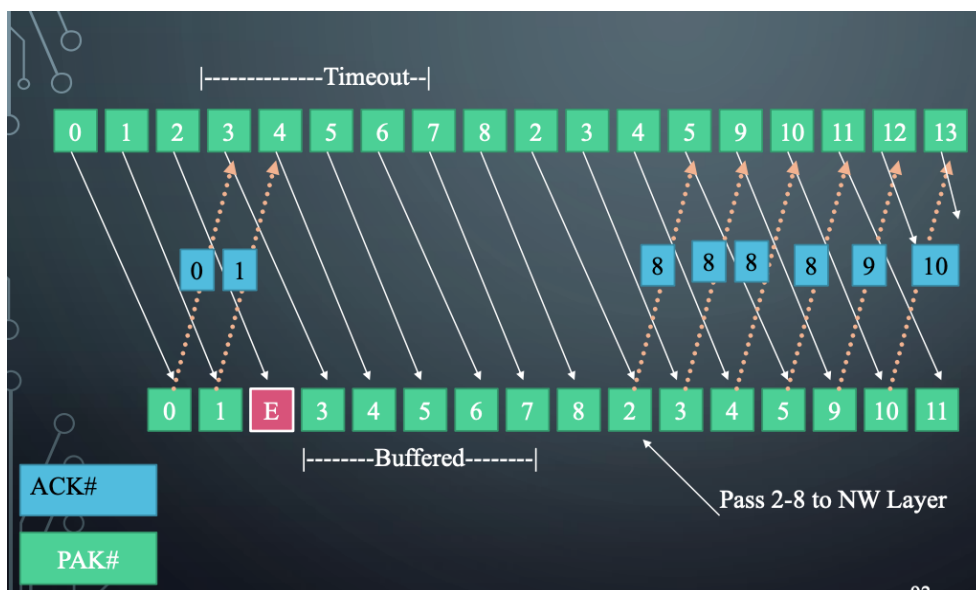


Fig. 3-12. A sliding window of size 1, with a 3-bit sequence number. (a) Initially. (b) After the first frame has been sent. (c) After the first frame has been received. (d) After the first acknowledgement has been received.

## Session With Recovery



### Stopping the Problem

- A Data-Link cannot be stopped.
- Consider a session termination.
- Neither terminal knows that other has sent last packet, the last packet must be ACK'd.
- In practice the data-link is dropped after the link is sensed as being dead for a prolonged period